Deep Learning Architectures for Rough Volatility Models: Mapping Rough Heston Parameters to Implied Volatility Surfaces

The Evolution of Volatility Modeling and the Rough Paradigm

The evolution of financial mathematics has been fundamentally driven by the continuous pursuit of models capable of accurately capturing the empirical realities of market dynamics. For decades, the foundational Black-Scholes framework, which assumed that underlying asset returns followed a geometric Brownian motion with constant volatility, served as the bedrock of option pricing theory. However, the assumption of constant volatility was empirically falsified by the market crash of 1987, after which equity markets began to exhibit a pronounced implied volatility smile and skew.[1] To rectify this inconsistency, the quantitative finance community transitioned toward local and stochastic volatility models. The most notable of these was the classical Heston model, introduced in 1993, which allowed variance to evolve dynamically as a mean-reverting square-root diffusion process.[3] While the classical Heston model successfully generated volatility smiles and accounted for the correlation between asset returns and volatility—known as the leverage effect—it inherently relied on the assumption that the stochastic volatility driver was a standard Brownian motion.[5]

This structural assumption of standard Brownian motion, which dictates that volatility paths have a Hölder regularity of 0.5, forced the term structure of the at-the-money (ATM) implied volatility skew to remain bounded and relatively flat as the time to maturity approached zero. This theoretical prediction directly contradicted empirical market data. In real option markets, the ATM volatility skew demonstrates a steep, explosive power-law behavior for short-dated options, reflecting the market's pricing of extreme tail risks and rapid microstructural reactions.[2]

A definitive paradigm shift occurred with the introduction of rough volatility models, pioneered by Gatheral, Jaisson, and Rosenbaum. Their seminal research demonstrated that the time series of historical realized variance is far rougher than standard Brownian motion.[5] By replacing the standard Brownian driver with a fractional Brownian motion (fBm) characterized by a Hurst exponent (H) strictly less than 0.5—typically clustered around 0.1 for broad market equity indices—rough volatility models successfully reproduced the stylized facts of high-frequency financial time series.[1] Under the risk-neutral pricing measure, the incorporation

of the fractional memory kernel allowed these models to perfectly replicate the $\tau^{H-1/2}$ power-law decay of the ATM volatility skew for small maturities ($\tau \to 0$), bridging the gap between historical volatility dynamics and the pricing of short-dated derivatives.[2]
However, the transition from classical stochastic volatility to rough fractional volatility introduced formidable computational bottlenecks that threatened the practical applicability of the models.[9] The lack of Markovianity and the presence of a singular memory kernel in the fractional Brownian motion fundamentally invalidated traditional pricing methods, such as standard finite difference schemes for partial differential equations (PDEs) and simple Monte Carlo simulations, which rely heavily on the independent increments of Markov processes.[3]
While the rough Heston model is uniquely amenable to characteristic function analysis via fractional Riccati equations, the numerical evaluation of these functions requires highly complex algorithms like the Adams predictor-corrector method or the SINH-acceleration Fourier inversion.[11] Consequently, calibrating the rough Heston model to a dense implied volatility surface requires thousands of evaluations of a highly expensive semi-analytical function. This transforms routine daily market calibrations into dense global optimization problems that can take hours to converge, thereby rendering the mathematical elegance of the model entirely impractical for real-time market-making, hedging, and algorithmic trading operations.[9]
To resolve this computational deadlock, the literature has increasingly turned to deep learning, utilizing artificial neural networks as ultra-fast, high-dimensional surrogate approximators. By learning the complex, non-linear mapping between rough volatility parameters and the implied volatility surface in an offline environment, deep neural networks effectively reduce the online calibration time from hours to milliseconds.[14] This report provides an exhaustive, comparative analysis of the deep learning architectures currently deployed for rough volatility calibration, progressing from foundational Multilayer Perceptrons (MLPs) to state-of-the-art physics-informed Residual Networks (ResNets), Graph Neural Operators (GNOs), and Fourier Neural Operators (FNOs).

The Mathematical Framework of the Rough Heston Model

To comprehend the exacting requirements placed upon any neural network tasked with performing deep calibration, it is essential to first understand the structural and mathematical complexity of the rough Heston model. The model specifies the joint dynamics of a financial asset price process S_t and its instantaneous variance process V_t under a risk-neutral pricing measure $\mathbb{Q}$. The system of stochastic differential equations (SDEs) governing these dynamics is formally defined as follows:

$$\frac{dS_t}{S_t} = \sqrt{V_t} dZ_t$$

$$V_t = V_0 + \int_0^t K(t-s)\lambda(\theta - V_s)ds + \int_0^t K(t-s)\nu\sqrt{V_s}dW_s$$

In this formulation, Z_t and W_t represent standard Brownian motions that are correlated by a coefficient $\rho \in [-1, 1]$. This correlation parameter is essential for capturing the leverage effect, dictating the asymmetry and downward slope of the equity implied volatility skew.[1] The parameter λ represents the speed of mean reversion, pushing the variance toward θ, which acts as the long-term mean variance level. The parameter ν (frequently denoted as η in associated literature) dictates the volatility of volatility, scaling the magnitude of variance fluctuations, while V_0 represents the initial forward variance state.[1] Furthermore, to ensure the weak existence of a solution and prevent the variance process from becoming negative, models must consider the Feller condition, which in classical formulations requires $2\lambda\theta \geq \nu^2$, though the fractional nature of the rough Heston model complicates strict boundary behaviors.[9]

The critical departure from the classical Heston model, and the source of its immense computational complexity, lies in the introduction of the memory kernel $K(u)$. In the rough Heston framework, the fractional kernel is explicitly defined as a Riemann-Liouville fractional integral kernel:

$$K(u) = \frac{u^{H-1/2}}{\Gamma(H + 1/2)}$$

Here, $H \in (0, 0.5)$ represents the Hurst exponent, which governs the roughness of the paths, and $\Gamma(\cdot)$ represents the mathematical Gamma function.[1] As the time difference $u \to 0$, the kernel $K(u)$ exhibits a severe singularity. This singularity is precisely the mathematical mechanism that generates the explosive, rough behavior of the volatility paths and allows the model to perfectly replicate the extreme skewness observed in options nearing expiration.[21]

Because the variance process V_t continuously relies on the entire historical path of the asset through the convolution integral, the resulting system is fundamentally non-Markovian.[1] Option

pricing under this framework requires computing the mathematical expectation of the discounted option payoff. For standard European options, this can be efficiently evaluated if the characteristic function of the log-price is analytically or semi-analytically known. El Euch and Rosenbaum demonstrated that the characteristic function of the rough Heston model can be elegantly expressed in terms of a solution to a fractional Riccati equation, which replaces the standard ordinary differential equation (ODE) found in the classical Heston model.[6]
Solving this fractional Riccati equation, however, involves complex numerical integration over the singular fractional kernel. Even when utilizing highly optimized numerical integration techniques, such as the Adams predictor-corrector scheme or advanced Fourier inversion techniques utilizing SINH-acceleration, generating a complete implied volatility surface for hundreds of distinct strike and maturity combinations remains highly computationally burdensome.[11] Furthermore, the objective function mapping the parameter space $\Theta = \{H, \nu, \rho, V_0, \lambda, \theta\}$ to the market's implied volatility surface is highly non-convex. The parameter landscape is riddled with local minima and characterized by flat gradients, particularly concerning the highly non-linear interaction between the Hurst exponent and the volatility of volatility.[10] This non-convexity necessitates the use of robust global optimization routines, which intrinsically require thousands of repeated evaluations of the pricing function, compounding the computational cost and cementing the necessity for deep neural network-based surrogates.[10]

Deep Calibration Frameworks: Forward Pricing versus Inverse Mapping

The application of deep learning algorithms to the rough volatility calibration problem is generally bifurcated into two distinct methodological approaches: the forward pricing approximation (often referred to as the two-step method) and the direct inverse mapping method. Both approaches seek to eliminate the bottleneck of evaluating the fractional Riccati equation during live trading.
In the forward pricing approximation approach, the neural network acts as a highly efficient, differentiable surrogate for the complex mathematical pricing function.[14] The network, denoted mathematically as $\mathcal{F}_\omega$ with a set of trainable synaptic weights ω, is trained to map the rough Heston parameter vector Θ and the specific contract specifications (such as the log-moneyness k and the time to maturity τ) to the corresponding implied volatility σ_{imp} or the raw option price. The training dataset for this network is generated entirely offline by simulating millions of parameter combinations drawn via Latin Hypercube Sampling from a predefined bounded hypercube $\Theta \in$. The precise "ground truth" implied volatilities for these parameters are computed using the slow, traditional fractional Riccati solvers.[14]
Once the neural network is sufficiently trained to a high degree of accuracy—often achieving a

Mean Squared Error (MSE) loss function below 10^{-4} on validation sets—it permanently replaces the traditional numerical solver in the online calibration phase.[13] The live calibration then proceeds using standard gradient-based numerical optimizers (such as the Levenberg-Marquardt algorithm or L-BFGS) to minimize the distance between the network's predicted implied volatility surface and the observed market surface:

$$\hat{\Theta} = \arg\min_{\Theta} \sum_{i,j} w_{i,j} \left(\sigma_{market}(k_i, \tau_j) - \mathcal{F}_\omega(\Theta, k_i, \tau_j)\right)^2$$

This two-step approach vastly benefits from the fact that evaluating the neural network and computing its gradients with respect to the input parameters via automatic differentiation is nearly instantaneous. This allows traditional, robust optimizers to navigate the non-convex landscape in milliseconds rather than hours, retaining the mathematical interpretability of the parameters while leveraging the speed of GPU-accelerated matrix multiplication.[14]

Conversely, the direct inverse mapping method attempts to train a neural network $\mathcal{G}_\phi$ to bypass the optimization step entirely, mapping the observed implied volatility surface directly back to the optimal rough Heston parameters.[12] In this network architecture, the entire implied volatility surface—typically formatted as a standardized spatial grid of shape $N_{strikes} \times M_{maturities}$ —is fed into the input layer of the network. The hidden layers extract the spatial relationships of the volatility smile, and the output layer yields the predicted parameter vector $\hat{\Theta}$.[10]

While this inverse approach theoretically reduces the online calibration time to a single forward pass through the network (achieving $\mathcal{O}(1)$ time complexity), it struggles significantly with out-of-sample generalization. Financial option markets are highly dynamic and notoriously noisy; derivative contracts constantly expire, new strikes are introduced based on underlying asset movements, and the spatial configuration of the valid implied volatility surface shifts continuously throughout the trading day.[26] Inverse mapping networks trained on a strictly fixed grid of inputs fail catastrophically when live market data is missing, sparse, or geometrically misaligned with the predefined training grid.[26] To utilize such models, quantitative analysts must rely on complex interpolation pre-processing to force the live data into the required grid shape, a practice that frequently introduces static arbitrage violations and invalidates the resulting parameter estimations.[26]

Standard Multilayer Perceptrons (MLPs) in Rough Volatility

The foundational deep learning architecture utilized to overcome the calibration bottleneck is the Multilayer Perceptron (MLP). In the early applications of deep learning to rough volatility, researchers such as Bayer, Stemper, Horvath, Muguruza, and Tomas demonstrated that properly configured MLPs could successfully approximate the mapping from the highly complex rough parameter space to the implied volatility surface.[30]

The typical MLP architecture engineered for this task involves flattening the two-dimensional implied volatility grid into a single one-dimensional input vector. For instance, an input surface comprising a grid of 8 maturities and 11 strikes is systematically flattened into an 88-dimensional vector. This vector is then processed sequentially through a series of fully connected dense hidden layers.[10] These hidden layers utilize non-linear activation functions to capture the complex, non-linear relationships inherent in the fractional Brownian motion dynamics.[10]

Crucially, the choice of activation function is not arbitrary. Advanced MLP implementations favor Exponential Linear Units (ELUs) or Softplus activations over standard Rectified Linear Units (ReLUs).[10] The choice of ELU is mathematically deliberate; ELU provides smooth, non-zero continuous gradients for negative inputs. This property is highly beneficial when the network must capture and backpropagate gradients related to the negative correlation parameter ($\rho \in [-1, 0]$), ensuring that the network's neurons do not "die" and gradients do not vanish during the rigorous optimization loops required by Levenberg-Marquardt calibration algorithms.[10]

A standard MLP utilized in rough Heston surface calibration might contain three to four fully connected hidden layers, cumulatively containing approximately 286,000 trainable parameters.[30] Despite their architectural simplicity relative to modern deep learning models, MLPs are mathematically proven universal approximators capable of delivering remarkable computational speedups.[24] Once trained offline, the MLP evaluation of an entire implied volatility surface requires only basic, highly parallelizable matrix multiplications. This shifts the calibration burden from numerical integration to simple floating-point operations, yielding calibration times in the realm of single-digit milliseconds compared to the hours required for standard Monte Carlo or numerical PDE methods.[14]

However, the rigid, fully connected nature of the MLP exposes several critical mathematical and practical vulnerabilities when applied specifically to the rough volatility regime.

Firstly, the MLP is inherently oblivious to the spatial and geometric relationships embedded within the implied volatility surface.[30] By flattening the 2D volatility grid into a 1D vector, the network actively discards the inductive bias that adjacent strikes and maturities are highly correlated. The network is forced to painstakingly learn these structural relationships from scratch through immense amounts of data, rendering the training process highly inefficient.[28]

Secondly, MLPs suffer profoundly from the curse of dimensionality and grid-locking.[35] If the network is trained to accept a specific 8×11 geometric grid, any live market data that deviates from this exact dimensional structure must be artificially interpolated, extrapolated, or discarded.[26] This severely limits the model's robustness in live trading environments, where liquidity varies wildly across the strike chain and the removal of noisy outliers is a standard quantitative practice.[26]

Finally, while MLPs perform adequately in capturing the broad, macroeconomic shape of the volatility smile, they frequently struggle to accurately resolve the $t \to 0$ mathematical singularity associated with the fractional Hurst exponent.[7] The extreme steepness of the short-maturity skew requires the neural network to navigate extraordinarily sharp gradients in the loss topology. Standard MLPs often suffer from approximation errors in their deeper layers as the optimization loss landscape becomes excessively flat, resulting in degraded calibration accuracy for the very short-dated options that rough volatility models were explicitly designed to price.[10]

Advanced Architectures: Deep Residual Networks (ResNets)

To mitigate the gradient degradation, grid rigidity, and structural approximation errors encountered by standard MLPs, researchers have increasingly transitioned toward utilizing Deep Residual Networks (ResNets) and Convolutional Neural Networks (CNNs) for the calibration of rough Heston and rough Bergomi models.[6] Originally developed by the computer vision community to solve the vanishing gradient problem in ultra-deep image classification networks, the ResNet architecture introduces skip (or residual) connections that allow mathematical information to bypass intermediate non-linear transformations.[10]

In the specific context of rough volatility calibration, a ResNet block modifies the standard feedforward transformation as follows:

$$\mathbf{h}^{(\ell+1)} = \sigma(\mathbf{W}^{(\ell)}\mathbf{h}^{(\ell)} + \mathbf{b}^{(\ell)}) + \mathbf{h}^{(\ell)}$$

where $\mathbf{h}^{(\ell)}$ represents the input vector to the ℓ-th layer, $\mathbf{W}$ and $\mathbf{b}$ are the weight matrices and bias vectors, and σ represents the non-linear activation function. By forcing the neural network to learn the residual mapping $\mathcal{F}(\mathbf{h}) = \mathbf{h}^{(\ell+1)} - \mathbf{h}^{(\ell)}$ rather than the direct mapping, ResNets provide an unimpeded gradient pathway back to the earliest layers of the network during the backpropagation phase of training.[10]

This seemingly simple architectural innovation holds profound mathematical implications for the calibration of the rough Heston model. The simultaneous calibration of the Hurst exponent (H) and the volatility of volatility (ν) is notoriously unstable due to their highly non-linear, compounding interaction near the origin of the temporal domain.[3] Extremely small perturbations in the Hurst parameter, particularly in the regime where $H \in (0.01, 0.15)$, cause massive, highly sensitive shifts in the short-dated implied volatility skew, creating a highly complex, rugged loss topology for the optimizer.[7] The skip connections inherent in ResNets ensure that the intricate, high-frequency features of the parameter-to-price mapping—such as the power-law explosion of the volatility skew—are strictly preserved and not filtered out by the successive smoothings of dense layers.[38]

Early implementations by researchers such as Henry Stone contextualized this problem by employing 1D and 2D Convolutional Neural Networks (CNNs) to directly ingest simulated fractional sample paths and map them to their corresponding Hölder exponents.[37] By utilizing convolutional filters, the network could intrinsically detect the rough structural patterns of the fractional Brownian motion, vastly outperforming MLPs in parameter estimation accuracy.[37] Expanding on this, modern calibration architectures incorporate Physics-Informed ResNets (PIRNs). These frameworks enforce the theoretical constraints of the rough Heston model, such as the fractional Feller condition and the boundaries of the Hurst parameter, directly into the network's architecture and composite loss function.[38] Furthermore, incorporating Batch Normalization layers between the residual blocks serves to stabilize the internal covariate shift. This ensures that the disparate numerical scales of the input parameters (for example, a highly sensitive initial variance $V_0 \approx 0.02$ versus a broader mean reversion speed $\lambda \approx 1.5$) do not destabilize the optimization dynamics.[10]

Empirical evidence decisively demonstrates that ResNet-based calibration architectures achieve substantially lower out-of-sample generalization errors than standard MLPs. They exhibit highly robust convergence behavior even when calibrating the singular boundaries of the rough volatility parameter space.[10] Despite these immense analytical advancements, ResNets and CNNs remain fundamentally constrained by the same fixed-grid input requirements that plague MLPs. Because they require standardized matrices to perform convolutions, they remain sub-optimal for processing the raw, unstructured point-cloud data generated by fragmented, live option markets.[26]

Graph Neural Networks (GNNs) and Operator Deep Smoothing

To systematically address the persistent rigidity of fixed-grid inputs, the deep calibration frontier has recently expanded to incorporate Graph Neural Networks (GNNs) and, more specifically, Graph Neural Operators (GNOs).[26] In live trading environments, financial option data does not naturally manifest as a dense, rectangular image grid. Instead, it arises dynamically in ever-changing, irregular spatial configurations dictated by available market liquidity, corporate actions, and the continuous decay of time to maturity.[26] When classical MLPs or ResNets attempt to process this raw data, any missing observations must be artificially interpolated, creating localized static arbitrage violations that irreversibly corrupt the subsequent calibration of sensitive rough volatility parameters.[26]

Graph Neural Networks conceptualize the implied volatility surface not as a rigid matrix, but as an unstructured geometric point cloud. In this representation, each node represents an individual traded option contract, uniquely defined by its continuous spatial coordinates (strike, maturity) and its observed volatility feature.[45] By employing complex message-passing algorithms, GNNs aggregate localized information from spatially neighboring nodes to construct a continuously defined mathematical surface without requiring uniform data spacing.[45]

The most significant theoretical leap in this domain is the development of the Graph Neural Operator (GNO), pioneered specifically for the task of "operator deep smoothing" of the implied volatility surface.[26] Unlike traditional neural networks that approximate a mapping between finite-dimensional Euclidean spaces, a neural operator is designed to learn a mapping between infinite-dimensional function spaces.[23] The GNO directly maps a dynamically sized input set of quoted market data $v(x_1), v(x_2), ..., v(x_n)$ evaluated at arbitrary points into a smooth, continuous implied volatility surface function $u(x)$ defined over the entire continuous domain of strikes and maturities.[26]

This discretization-invariant property of the GNO means that a single, pre-trained network instance can accurately ingest an S&P 500 option chain containing 500 valid quotes on a highly liquid Monday, and subsequently ingest 450 entirely different quotes on a less liquid Tuesday, without requiring any architectural modification or hazardous data pre-processing.[26] The integration of the GNO framework entirely bypasses the instance-by-instance parameter optimization that hinders traditional parametric implied volatility approaches like SVI (Stochastic Volatility Inspired) and SSVI.[26]

Moreover, advanced variations of the operator approach, such as the HyperIV architecture, leverage the concept of hypernetworks to augment GNOs for high-frequency applications.[49] In this state-of-the-art framework, a master network instantaneously processes the unstructured market data point cloud to dynamically generate the weights for a smaller, highly compact secondary neural network. This secondary network mathematically defines the complete implied volatility surface.[49]

Crucially, by directly incorporating hard penalization terms for calendar spread arbitrage and butterfly spread arbitrage into the composite loss function during the offline training phase, the resulting GNO ensures that the inferred rough volatility surface is strictly consistent with the theoretical constraints of no-arbitrage pricing theory.[26] The HyperIV architecture demonstrates superior real-time performance, generating high-quality implied volatility surfaces in approximately 2 milliseconds using as few as 9 observed market option prices. It significantly outperforms standard models, Variational Autoencoders (VAEs), and basic MLPs in both computational speed and predictive accuracy, cementing its utility in fast-paced trading environments.[49]

Fourier Neural Operators (FNOs) and the Fractional Riccati Equation

While Graph Neural Operators excel at the spatial interpolation and smoothing of the observed implied volatility surface, Fourier Neural Operators (FNOs) represent the absolute vanguard of solving the underlying stochastic partial differential equations (SPDEs) and path-dependent PDEs (PPDEs) that govern the rough Heston model.[6] The valuation of derivatives under rough volatility corresponds to finding the solution to a highly complex PPDE driven by fractional Brownian motion.[9] Standard numerical solvers for these equations are plagued by the curse of dimensionality and the severe mathematical singularity located along the diagonal of the fractional integration kernel.[9]

The Fourier Neural Operator circumvents this numerical catastrophe by parameterizing the integral operator entirely in the frequency domain.[23] By applying the Fast Fourier Transform (FFT), the FNO filters out high-frequency numerical noise and learns the dominant spectral modes of the fractional Riccati equation. It then applies an inverse FFT to return the computed solution to the spatial domain.[53] This powerful spectral inductive bias enables the FNO to learn the highly non-linear operator mapping the parameter space Θ to the infinite-dimensional solution space of the PPDE with near-linear time complexity. The FNO evaluates significantly faster than conditional diffusion generative models or massive Monte Carlo rollouts.[23]

However, the direct application of classical FNOs to the rough volatility problem encounters a profound mathematical obstacle: the inherent periodicity assumption of the Fast Fourier Transform.[56] Fractional Brownian motion paths, by their mathematical definition on a finite trading interval $$, are strictly non-periodic.[57] Forcing non-periodic fractional stochastic paths through a standard FNO introduces severe boundary artifacts and triggers the Gibbs phenomenon. This spectral leakage deeply corrupts the valuation of the options and wildly distorts the calibration of the Hurst exponent.[56]

To resolve this critical limitation, the computational finance literature introduced the Mirror-Padded Fourier Neural Operator (MFNO).[57] The MFNO architecture mathematically extends the fractional stochastic path by reflecting it symmetrically about the terminal midpoint $t = T$, thus extending the temporal domain from

$$to$$

.[56] The function values on the extended domain $(T, 2T]$ become the exact mirror image of the original historical path, yielding a mathematically continuous and strictly periodic function over the entire $$ interval.[56]

Rigorous proofs utilizing Wong-Zakai type approximation theorems demonstrate that MFNOs can approximate the solutions of path-dependent SDEs and Lipschitz transformations of fractional Brownian motions to an arbitrary degree of accuracy, entirely bypassing the boundary limitations of standard FNOs.[57] Crucially, the MFNO exhibits robust zero-shot resolution generalization. This property allows the operator, which was trained on a coarse temporal grid to save computational resources, to evaluate option prices on highly dense temporal grids without requiring any retraining or fine-tuning.[57]

Furthermore, standard neural operators typically collapse complex stochastic processes down to their conditional mean, discarding the critical variance and tail structures that are absolutely essential for accurate option pricing under rough volatility.[23] To counter this, researchers have developed the Martingale Neural Operator (MNO), which incorporates a deep architectural prior based on the Doob-Meyer decomposition theorem.[23] The Doob-Meyer theorem dictates that any submartingale can be uniquely decomposed into a predictable, finite-variation drift component and a zero-mean martingale component.[23] The MNO embeds this factorization directly into its neural architecture, mapping the initial market conditions directly to both the conditional mean and the covariance of the terminal law.[23] Evaluated extensively on rough volatility models with Hurst parameters as low as $H = 0.1$, the MNO accurately preserves the heavy distributional tails and reduces the Wasserstein distance to the true distribution by immense magnitudes (up to $120\times$) compared to sequential baselines like LSTMs or Neural CDEs.[23]

Comparative Architectural Analysis

The deployment of these highly diverse deep learning architectures highlights distinct tradeoffs regarding computational cost, mathematical inductive bias, and spatial flexibility when tackling the rough Heston model. The following structured analysis encapsulates the comparative advantages and limitations of each framework currently utilized in the literature:

Architecture	**Core Mathematical Mechanism**	**Fixed Input Required?**	**Optimization & Generalizatio n Profile**	**Primary Application in Rough Volatility**
Multilayer Perceptron (MLP)	Dense affine transformation s with non-linear continuous activations (ELU/Softplus).	Yes. Requires rigid $N \times M$ 2D grids mathematically flattened to 1D vectors.	High risk of over-fitting and grid-locking. Frequently struggles with the highly singular $t \to 0$ boundary of the Hurst exponent.	Baseline fast surrogate pricer. Operates well only in highly standardized, deeply liquid market conditions.
Residual Network (ResNet)	Skip connections propagating the identity function ($\mathbf{h}^{(\ell+1)} = \mathcal{F}(\mathbf{h}^{($	Yes. Operates on standardized image-like matrices and tensors.	Mitigates vanishing gradients. Smooths the highly non-convex parameter	Physics-Inform ed mapping; highly robust calibration of extremely volatile

	).		landscape, significantly improving test-set accuracy over MLPs.	parameters (ν and H).
Graph Neural Operator (GNO)	Message passing across nodes in infinite-dimensional functional spaces.	No. Fully discretization-invariant. Directly ingests raw, unstructured point clouds.	Excellent robustness against missing data. Enforces static no-arbitrage constraints globally rather than point-wise.	"Operator Deep Smoothing" and real-time, high-frequency implied volatility surface nowcasting.
Fourier Neural Operator (FNO / MFNO)	Spectral convolutions utilizing Fast Fourier Transforms (FFTs). Mirror-padding enforces strict boundary periodicity.	No. Resolution-invariant. Evaluates continuous mathematical functions.	Zero-shot resolution generalization. Preserves the Wong-Zakai convergence theorems for non-Markovian fractional paths.	Directly solving the fractional Riccati PPDEs; simulation of non-Markovian fractional Brownian motion paths.
Martingale Neural Operator (MNO)	Embeds the Doob-Meyer decomposition into the neural graph via low-rank positive semi-definite factorization.	No. Maps initial conditions to terminal law covariance matrices.	Vastly reduces Wasserstein distance to the true distribution. Explicitly preserves stochastic variance and tail risk over conditional means.	Pricing complex exotics and capturing extreme tail distributions in deeply rough regimes ($H \approx 0.1$).

Second and Third-Order Market Insights

The transition from classical numerical calibration schemes to advanced operator-based deep learning yields profound implications that extend far beyond raw computational speedups. Synthesizing the underlying mathematical and empirical data reveals several crucial second and third-order insights regarding the future trajectory of quantitative volatility modeling.
Firstly, the overwhelming success of Graph Neural Operators and Hypernetworks in performing "discretization-invariant" mapping signals a profound philosophical shift from local parameter fitting to global market learning. Traditional rough Heston calibration is an instance-by-instance, highly localized process: a quantitative analyst takes a single snapshot of the market, runs an iterative numerical optimizer, and extracts the parameters valid only for that specific, isolated moment.[26] Conversely, a properly structured GNO trained over millions of historical market paths has effectively mapped the entirety of the market's dynamic, macro-level topology.[26] This establishes the foundation for entirely "hands-off" volatility modeling, wherein the neural operator continuously ingests unstructured raw tick data and dynamically outputs an arbitrage-free implied volatility surface without ever requiring a localized numerical recalibration or human intervention.[26]
Secondly, the integration of Doob-Meyer factorizations in Martingale Neural Operators reveals a critical, rapidly accelerating convergence between pure mathematical finance and deep learning architecture design.[23] The historical failure of standard neural networks to capture stochastic variance and tail risk represents a fundamental incompatibility between deterministic function approximation and the realities of stochastic calculus.[23] By embedding the theoretical decomposition of martingales and predictable drifts directly into the neural network's computational graph, the network ceases to be a mere "black box" statistical approximator. Instead, it becomes a structural mathematical entity governed directly by the rules of Itô calculus.[9] This alignment suggests that the ultimate solution to the rough volatility bottleneck is not simply increasing the depth, width, or parameter count of an MLP, but structurally constraining the network using the governing stochastic differential equations of the financial asset class.[42]
Furthermore, the empirical observation that universal, parameter-free deep learning models (such as broad LSTMs or global GNOs) perform on par with tightly constrained parametric rough volatility models (where $H \approx 0.1$) provides powerful nonparametric evidence of a universal volatility formation process across asset classes.[8] The fact that both the explicit rough fractional stochastic volatility framework and entirely model-free neural architectures converge on the exact same predictive accuracy suggests that the fractional scaling behavior observed in the market is an intrinsic structural property of liquidity and order flow.[8] This validates the theoretical premise of the rough volatility paradigm independent of the specific mathematical formulation used to express it, cementing rough volatility not as a mathematical curiosity, but as a fundamental law of modern market microstructure.

Concluding Analysis

The calibration of rough volatility models, specifically the rough Heston architecture, represents one of the most computationally demanding and mathematically intricate challenges in modern quantitative finance. The singular fractional memory kernel that allows these models to perfectly capture the short-maturity dynamics of the implied volatility skew inherently breaks the Markovian assumptions that underpin traditional numerical PDE solvers. This structural reality renders classical calibration techniques entirely obsolete for high-frequency, real-time trading applications.

The integration of deep learning architectures provides a definitive and highly scalable resolution to this bottleneck. While standard Multilayer Perceptrons (MLPs) offer foundational computational speedups by treating calibration as a direct inverse mapping or surrogate pricing problem, they remain heavily constrained by geometric grid rigidities and severe gradient degradation near the Hurst parameter's boundary singularity. The evolution toward Deep Residual Networks (ResNets) introduces necessary structural stability through skip connections and batch normalization, allowing for the deeper approximation of the highly non-convex calibration landscape and the robust identification of negative correlations. However, the absolute vanguard of the field lies in the adoption of operator learning. Graph Neural Operators fundamentally resolve the issue of unstructured, dynamically shifting market data, executing discretization-invariant deep smoothing to guarantee arbitrage-free volatility surfaces irrespective of missing or noisy market quotes. Simultaneously, advanced Fourier Neural Operators, explicitly augmented with mirror-padding to handle non-periodic paths and martingale factorizations to capture tail risk, bridge the historical gap between deterministic deep learning and fractional stochastic calculus. By embedding the mathematical realities of non-periodic fractional Brownian motion and the Doob-Meyer decomposition directly into the network architecture, these advanced operators offer arbitrary accuracy in solving path-dependent PDEs with powerful zero-shot resolution generalizability. Ultimately, these deep learning frameworks transition rough volatility modeling from a mathematically elegant but computationally intractable theory into a robust, real-time mechanism capable of dominating modern derivatives pricing and risk management.

Works cited

1. From Stochastic to Rough Volatility: A New Deep Learning Perspective on Hedging - MDPI, accessed May 27, 2026, https://www.mdpi.com/2504-3110/7/3/225
2. On short-time behavior of implied volatility in a market model with indexes 1 - arXiv, accessed May 27, 2026, https://arxiv.org/html/2402.16509v2
3. A neural network rough volatility model - ResearchGate, accessed May 27, 2026, https://www.researchgate.net/profile/Jingtang-Ma/publication/373141528_A_neural_network_rough_volatility_model/links/64dc3071ad846e2882946115/A-neural-network-rough-volatility-model.pdf
4. Simulations of Implied Volatility and Option Pricing using Neural Networks and

Finite Difference Methods for Heston Model - Nazarbayev University, accessed May 27, 2026, https://nur.nu.edu.kz/bitstreams/4019723e-4633-44ff-8d22-fbfca0910dc4/download
5. Tesi di Laurea Financial Modeling under Rough and Fractional Stochastic Volatility, accessed May 27, 2026, https://thesis.unipd.it/retrieve/03223182-1c46-47b8-903b-3106832b6fe9/Dapporto_Antonio.pdf
6. Random Neural Networks for Rough Volatility - PMC - NIH, accessed May 27, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC12967488/
7. Impact of rough stochastic volatility models on long-term life insurance pricing - PMC, accessed May 27, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC9243767/
8. On the universality of the volatility formation process: When machine learning and rough volatility agree - American Institute of Mathematical Sciences, accessed May 27, 2026, https://www.aimsciences.org/article/doi/10.3934/fmf.2024002
9. Deep PPDES FOR ROUGH (LOCAL) STOCHASTIC VOLATILITY - Imperial College London, accessed May 27, 2026, https://www.ma.imperial.ac.uk/~ajacquie/index_files/JO%20-%20RoughVolPPDE_Final.pdf
10. Sunlight Analysis in Heston Model Calibration | PDF | Volatility (Finance) | Deep Learning, accessed May 27, 2026, https://www.scribd.com/document/871468262/Heston-Model-Calibration-via-Deep-Neural-Networks-1748284226
11. Correct implied volatility shapes and reliable pricing in the rough Heston model - arXiv, accessed May 27, 2026, https://arxiv.org/abs/2412.16067
12. Deep calibration with random grids - arXiv, accessed May 27, 2026, https://arxiv.org/html/2306.11061v2
13. Fast reliable pricing and calibration of the rough Heston model - arXiv, accessed May 27, 2026, https://arxiv.org/html/2508.15080
14. Deep Learning Volatility - King's Research Portal, accessed May 27, 2026, https://kclpure.kcl.ac.uk/portal/files/139539661/DLV_Revision.pdf
15. [1901.09647] Deep Learning Volatility - arXiv, accessed May 27, 2026, https://arxiv.org/abs/1901.09647
16. a deep neural network perspective on pricing and calibration in (rough) volatility models - Technical University of Munich, accessed May 27, 2026, https://portal.fis.tum.de/en/publications/deep-learning-volatility-a-deep-neural-network-perspective-on-pri/
17. On Deep Calibration of (rough) Stochastic Volatility Models | The Journal of FinTech, accessed May 27, 2026, https://www.worldscientific.com/doi/10.1142/S2705109925500051
18. Rough volatility: An overview - Baruch MFE Program, accessed May 27, 2026, https://mfe.baruch.cuny.edu/wp-content/uploads/2018/02/RoughVolatilityColumbia2018.pdf
19. Calibrating the Hurst Parameter for Rough Volatility Models with Application in the

South African Market - University of Cape Town, accessed May 27, 2026, https://open.uct.ac.za/bitstream/11427/37757/1/thesis_com_2022_pettit%20paul.pdf
20. Calibration and pricing under modified Heston non-affine stochastic volatility models - POLITesi, accessed May 27, 2026, https://www.politesi.polimi.it/retrieve/241853c3-63d7-458d-b4c0-d5ae5a02149c/MANNA%20Angelo%20CS%20-%20Tesi%20LM%20-%20version%202022-09-12.pdf
21. Deep learning interpretability for rough volatility - ResearchGate, accessed May 27, 2026, https://www.researchgate.net/publication/386336134_Deep_learning_interpretability_for_rough_volatility
22. Deep learning interpretability for rough volatility - arXiv, accessed May 27, 2026, https://arxiv.org/html/2411.19317v1
23. Martingale Neural Operators: Learning Stochastic Marginals via Doob-Meyer Factorization, accessed May 27, 2026, https://arxiv.org/html/2605.15806v1
24. A neural network-based framework for financial model calibration - arXiv, accessed May 27, 2026, https://arxiv.org/pdf/1904.10523
25. (PDF) Residual Expansion Approach for a Deep Learning Based Inverse Mapping Model Calibration - ResearchGate, accessed May 27, 2026, https://www.researchgate.net/publication/397542622_Residual_Expansion_Approach_for_a_Deep_Learning_Based_Inverse_Mapping_Model_Calibration
26. Operator Deep Smoothing for Implied Volatility - arXiv, accessed May 27, 2026, https://arxiv.org/html/2406.11520v3
27. Operator Deep Smoothing for Implied Volatility - arXiv, accessed May 27, 2026, https://arxiv.org/html/2406.11520v1
28. Operator Deep Smoothing for Implied Volatility - arXiv, accessed May 27, 2026, https://arxiv.org/html/2406.11520v2
29. Deep Smoothing of the Implied Volatility Surface | Request PDF - ResearchGate, accessed May 27, 2026, https://www.researchgate.net/publication/334033879_Deep_Smoothing_of_the_Implied_Volatility_Surface
30. Volatility Surface Reconstruction using Deep Learning under No-Arbitrage Constraints - arXiv, accessed May 27, 2026, https://arxiv.org/pdf/2605.24031
31. Volatility Model Calibration With Convolutional Neural Networks | Request PDF, accessed May 27, 2026, https://www.researchgate.net/publication/328302236_Volatility_Model_Calibration_With_Convolutional_Neural_Networks
32. Volatility Surface Reconstruction Using Deep Learning Under No-Arbitrage Constraints, accessed May 27, 2026, https://arxiv.org/html/2605.24031v1
33. Efficient Calibration for Option Pricing via a Physics-Informed Chebyshev Kolmogorov-Arnold Network - Preprints.org, accessed May 27, 2026, https://www.preprints.org/manuscript/202603.0678
34. Deep Learning vs. Black-Scholes: Option Pricing Performance on Brazilian Petrobras Stocks, accessed May 27, 2026, https://arxiv.org/html/2504.20088v1

35. (PDF) A neural network-based framework for financial model calibration - ResearchGate, accessed May 27, 2026, https://www.researchgate.net/publication/335672074_A_neural_network-based_framework_for_financial_model_calibration
36. Pricing and calibration of stochastic models via neural networks - Imperial College London, accessed May 27, 2026, https://www.imperial.ac.uk/media/imperial-college/faculty-of-natural-sciences/department-of-mathematics/math-finance/TOMAS_MEHDI_01390785.pdf
37. Calibrating rough volatility models: a convolutional neural network approach - IDEAS/RePEc, accessed May 27, 2026, https://ideas.repec.org/p/arx/papers/1812.05315.html
38. Adaptive Movement Sampling Physics-Informed Residual Network (AM-PIRN) for Solving Nonlinear Option Pricing models - arXiv, accessed May 27, 2026, https://arxiv.org/html/2504.03244v1
39. Prediction-Enhanced Monte Carlo: A Machine Learning View on Control Variate - arXiv, accessed May 27, 2026, https://arxiv.org/pdf/2412.11257?
40. Machine learning for option pricing: an empirical investigation of network architectures, accessed May 27, 2026, https://www.researchgate.net/publication/372416564_Machine_learning_for_option_pricing_an_empirical_investigation_of_network_architectures
41. Calibrating rough volatility models: a convolutional neural network approach, accessed May 27, 2026, https://www.researchgate.net/publication/335433816_Calibrating_rough_volatility_models_a_convolutional_neural_network_approach
42. A Hybrid PINNs-FNO Approach for Advanced Option Pricing Models - Preprints.org, accessed May 27, 2026, https://www.preprints.org/manuscript/202501.0629
43. Operator Deep Smoothing for Implied Volatility - IDEAS/RePEc, accessed May 27, 2026, https://ideas.repec.org/p/arx/papers/2406.11520.html
44. [2406.11520] Operator Deep Smoothing for Implied Volatility - arXiv, accessed May 27, 2026, https://arxiv.org/abs/2406.11520
45. Heterogeneous Graph Attention Networks for Realized Volatility Forecasting - kth .diva, accessed May 27, 2026, https://kth.diva-portal.org/smash/get/diva2:1987620/FULLTEXT01.pdf
46. SpotV2Net: Multivariate Intraday Spot Volatility Forecasting via Vol-of-Vol-Informed Graph Attention Networks - arXiv, accessed May 27, 2026, https://arxiv.org/html/2401.06249v1
47. Graph-Based Stock Volatility Forecasting with Effective Transfer Entropy and Hurst-Based Regime Adaptation - MDPI, accessed May 27, 2026, https://www.mdpi.com/2504-3110/9/6/339
48. Operator Deep Smoothing for Implied Volatility - OpenReview, accessed May 27, 2026, https://openreview.net/forum?id=DPlUWG4WMw
49. HyperIV: Real-time Implied Volatility Smoothing - OpenReview, accessed May 27, 2026, https://openreview.net/forum?id=o9jtS7HupI
50. ICML Poster HyperIV: Real-time Implied Volatility Smoothing, accessed May 27,

2026, https://icml.cc/virtual/2025/poster/44077
51. HyperIV: Real-time implied volatility smoothing - Edinburgh Research Explorer, accessed May 27, 2026, https://www.research.ed.ac.uk/files/526427109/YangEtalICML2025HyperIVReal-time.pdf
52. Deep Curve-Dependent PDEs for Affine Rough Volatility | SIAM Journal on Financial Mathematics, accessed May 27, 2026, https://epubs.siam.org/doi/10.1137/19M1267805
53. Forcing and Diagnosing Failure Modes of Fourier Neural Operators Across Diverse PDE Families - arXiv, accessed May 27, 2026, https://arxiv.org/html/2601.11428v1
54. Domain Agnostic Fourier Neural Operators | Lehigh Preserve, accessed May 27, 2026, https://preserve.lehigh.edu/lehigh-scholarship/faculty-and-staff-publications/faculty-publications/domain-agnostic-fourier
55. Fast Pricing and Greeks Computation for Barrier Options via Physics, accessed May 27, 2026, https://thesis.unipd.it/retrieve/9e9db9a6-1d81-4005-a07b-263e3a32b140/Tiboni_Gabriele.pdf
56. Fourier Neural Operators for Non-Markovian Processes:Approximation Theorems and Experiments - arXiv, accessed May 27, 2026, https://arxiv.org/pdf/2507.17887
57. Fourier Neural Operators for Non-Markovian Processes: Approximation Theorems and Experiments - arXiv, accessed May 27, 2026, https://arxiv.org/html/2507.17887v1
58. How a Fourier Neural Operator Learns to Solve PDEs — and Where It Falls Short - PhysicsX, accessed May 27, 2026, https://www.physicsx.ai/newsroom/how-a-fourier-neural-operator-learns-to-solve-pdes----and-where-it-falls-short
59. Volatility is rough | Request PDF - ResearchGate, accessed May 27, 2026, https://www.researchgate.net/publication/323818641_Volatility_is_rough
60. Research(상세) | Whats New | News - KIAS(Korea Institute For Advanced Study), accessed May 27, 2026, https://www.kias.re.kr/kias/bbs/B0000047/view.do?nttId=6665&menuNo=408048
61. Computational Methods in Financial Mathematics: Mirror-Padded Fourier Neural Operators and Penalized Backward Stochastic Differential Equations - 서울대학교 수리과학부, accessed May 27, 2026, https://www.math.snu.ac.kr/bbs/board.php?bo_table=Math_Seminar&wr_id=45715
62. Fourier Neural Operators for Non-Markovian Processes: Approximation Theorems and Experiments | Request PDF - ResearchGate, accessed May 27, 2026, https://www.researchgate.net/publication/403534286_Fourier_Neural_Operators_for_Non-Markovian_Processes_Approximation_Theorems_and_Experiments
63. [2507.17887] Fourier Neural Operators for Non-Markovian Processes:Approximation Theorems and Experiments - arXiv, accessed May 27, 2026, https://arxiv.org/abs/2507.17887
64. Multivariate Realized Volatility Forecasting with Graph Neural Network -

IDEAS/RePEc, accessed May 27, 2026,
https://ideas.repec.org/p/arx/papers/2112.09015.html