Comprehensive Analysis of Numerical Simulation Schemes for the Rough Heston Model: Asymptotic Complexity, Stability, and GPU Acceleration

The Paradigm Shift to Rough Volatility Models

The transition from classical stochastic volatility frameworks to the rough volatility paradigm constitutes one of the most critical theoretical and operational shifts in modern quantitative finance. For decades, continuous-time models, epitomized by the standard Heston model, relied on the foundational assumption that the instantaneous volatility of a financial asset is driven by a standard Brownian motion. While mathematically tractable, yielding semimartingale structures and admitting closed-form characteristic functions through classical affine Volterra equations, standard Brownian drivers consistently fail to reproduce a multitude of critical stylized facts observed in empirical market data. Chief among these empirical shortcomings is the model's inability to accurately reproduce the power-law explosion of the implied volatility skew for short-dated options, a phenomenon uniformly documented across global equity and foreign exchange markets.[1]

Extensive empirical investigations into high-frequency realized variance time series have definitively demonstrated that the log-volatility of financial assets behaves in a manner statistically indistinguishable from a fractional Brownian motion (fBm) characterized by a Hurst parameter H significantly less than $1/2$. Observations consistently yield estimates of $H \approx 0.1$ across diverse asset classes.[2] This low Hurst index fundamentally violates the standard Brownian assumption (where $H = 1/2$). Instead, it implies a "rough" volatility path that is far less smooth, exhibiting strong short-term memory, heavy-tailed increments, and a highly realistic representation of market microstructure dynamics. From a microstructural perspective, the macroscopic rough volatility models can be rigorously derived as the scaling limits of bivariate, cumulative, heavy-tailed Hawkes processes, establishing a profound link between high-frequency order book dynamics and macroscopic option pricing.[3]

The Rough Heston model was introduced to preserve the fundamental structural advantages of the classical Heston model—namely its affine nature—while capturing the empirical reality of rough volatility. This is achieved by substituting the standard Brownian driver of the variance process with a fractional Volterra process. Within this framework, the instantaneous variance process V_t is mathematically expressed as a stochastic Volterra integral equation incorporating a singular fractional kernel. Specifically, the process takes the form:

$$V_t = V_0 + \frac{1}{\Gamma(\alpha)} \int_0^t (t-s)^{\alpha-1} \lambda(\theta - V_s) ds + \frac{1}{\Gamma(\alpha)} \int_0^t (t-s)^{\alpha-1} \nu \sqrt{V_s} dW_s$$

,

where $\alpha = H + 1/2$, λ denotes the mean reversion speed, θ represents the long-term mean variance, and ν stands for the volatility of volatility.[6]

Despite maintaining an affine structure, which theoretically allows the conditional characteristic function of the log-price to be computed, the introduction of the fractional kernel $(t-s)^{\alpha-1}$ fundamentally destroys both the Markovian and semimartingale properties of the underlying stochastic process.[1] The destruction of the Markovian property induces immense computational complexities. Classical partial differential equation (PDE) methods, traditionally deployed via Feynman-Kac theorems, cannot be directly applied to non-Markovian systems. Similarly, standard Monte Carlo simulation methodologies suffer from severe complexity overhead due to the necessity of integrating over the entire historical trajectory of the process at every discrete time step. Consequently, devising highly efficient, stable, and scalable numerical schemes to solve the fractional equations for Fourier-based pricing, as well as to simulate the non-Markovian paths for Monte Carlo valuation, has emerged as a paramount challenge within the computational finance community.[2]

This report presents an exhaustive comparative analysis of the principal numerical methodologies developed to operationalize the rough Heston model. The evaluation focuses on the fractional Adams-Bashforth-Moulton method, Hybrid numerical schemes (including the Hybrid Quadratic Exponential algorithm), and Markovian approximations (such as the Lifted Heston model). The analysis explicitly compares these paradigms regarding their theoretical asymptotic complexity, their numerical stability in hyper-rough regimes, and their architectural suitability for massively parallel execution on Graphics Processing Units (GPUs).

Fourier Pricing and the Fractional Riccati Equation

In standard affine stochastic volatility models, pricing European options is typically executed in the frequency domain. The characteristic function of the asset's log-price is derived, and the option price is subsequently recovered via Fourier inversion techniques, such as the Carr-Madan formula or the Lewis formula. The rough Heston model preserves this affine structure, meaning that the characteristic function of the log-price can theoretically be expressed in terms of the solution to a Riccati equation.[8]

However, the transition from classical to rough volatility radically alters the nature of this equation. In the classical Heston model, the characteristic function is governed by a standard ordinary differential equation (ODE) of the Riccati type, which admits an explicit, closed-form solution. Conversely, in the rough Heston model, the presence of the fractional integration kernel transforms the governing equation into a nonlinear parametric fractional Riccati differential equation.[8]

The fractional Riccati equation takes the generalized form $D^\alpha h(a,t) = F(a, h(a,t))$, where D^α denotes the fractional derivative operator of order $\alpha \in (1/2, 1)$, and F represents a nonlinear function involving the model parameters and the complex variable a.[11] Crucially, this nonlinear parametric fractional Riccati differential equation does not admit a closed-form analytical solution under any parameter regime. The absence of a closed-form expression dictates that any calibration or pricing algorithm utilizing the Fourier inversion approach must rely entirely on numerical discretization schemes to evaluate the characteristic function at every required frequency node.[8] Given that robust model calibration often requires evaluating the characteristic function millions of times across various strikes, maturities, and optimization iterations, the numerical efficiency and stability of the chosen fractional Riccati solver become the primary determinants of the model's practical viability.[8]

The Fractional Adams-Bashforth-Moulton Method

The most direct and historically prevalent mathematical approach to solving the fractional Riccati equation associated with the rough Heston model is the fractional Adams-Bashforth-Moulton method. Often referred to simply as the fractional Adams method, this technique is an advanced numerical adaptation of classical predictor-corrector algorithms conventionally utilized for standard ODEs, specifically extended to accommodate fractional derivatives defined in the Riemann-Liouville or Caputo sense.[6]

Algorithmic Mechanics and the Historical Convolution

The fractional Adams algorithm approaches the fractional Riccati equation by discretizing the continuous time domain into a uniform grid of N discrete steps, characterized by a constant spacing parameter Δ. Because the fractional derivative is inherently a non-local mathematical operator, it possesses infinite memory. Consequently, calculating the value of the solution function $h(a, x)$ at the current step j is not solely dependent on the immediately preceding state $j-1$. Instead, it explicitly requires a discrete convolution calculated over all previously computed historical values from step 0 through step $j-1$.[13]

The methodology proceeds in a two-stage iterative loop for each time step. The initial stage is the predictor step (the Adams-Bashforth component), which computes an intermediate explicit approximation of the solution. This predicted value is generated using a weighted summation of all historical evaluations of the nonlinear function F. Following the predictor stage, the algorithm executes the corrector step (the Adams-Moulton component). The corrector refines the intermediate approximation utilizing an implicit formulation that incorporates the newly predicted forward value, subject to a distinct set of fractional integration weights. The fundamental operational characteristic—and primary vulnerability—of this scheme is its absolute reliance on historical heredity. To advance the numerical solution from time t_k to t_{k+1} , the computational engine must systematically retrieve from memory and process the entire history of the process trajectory $(t_0, t_1, \ldots, t_k)$.[13]

Asymptotic Time Complexity of Rough Heston Numerical Schemes

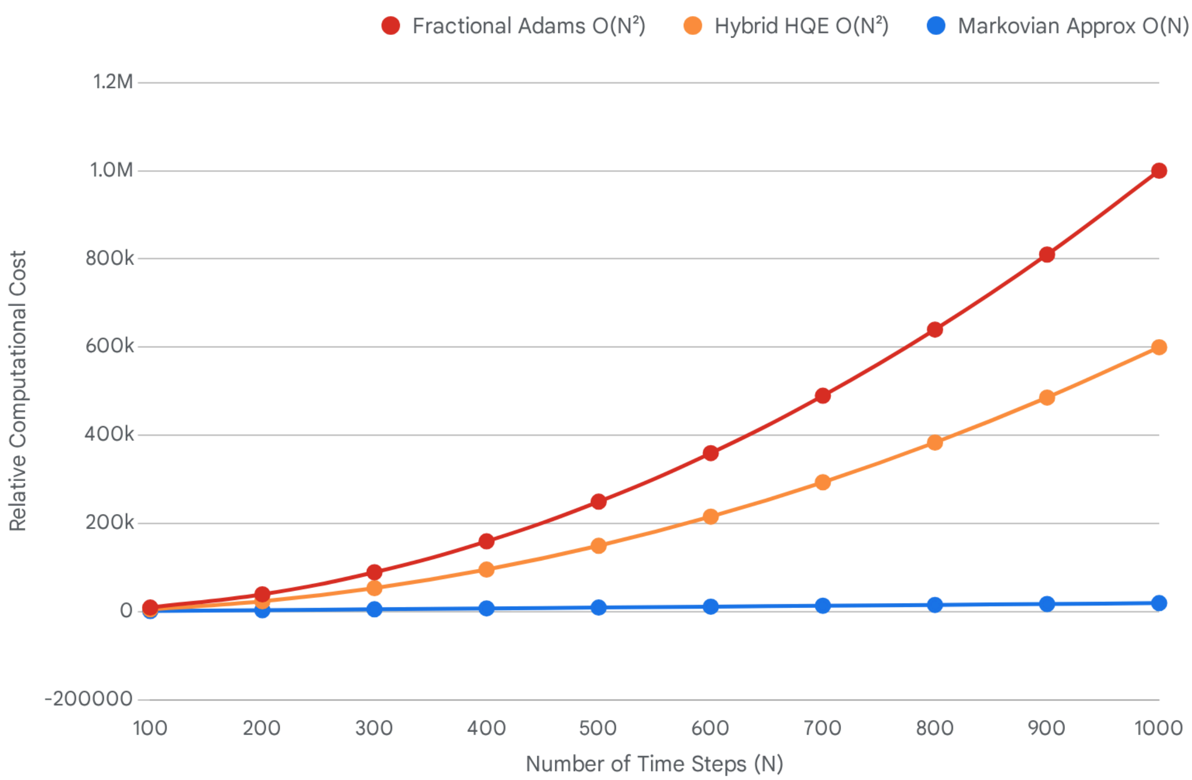

The Fractional Adams and pure HQE methods exhibit quadratic $O(N^2)$ scaling due to the need to compute stochastic Volterra integrals over the entire path history. In contrast, the Markovian Approximation (Lifted Heston) scales linearly O(N), enabling significantly faster computation for large numbers of time steps.

Data sources: Quantitative Finance, arXiv, MDPI, UPM

Asymptotic Computational Complexity

The mathematical requirement to iteratively compute a discrete-time convolution over a steadily expanding historical array at every forward step imposes an exceptionally severe computational penalty. Consequently, the core algorithmic complexity required to evaluate a single option price using the standard fractional Adams predictor-corrector method scales asymptotically as $\mathcal{O}(N^2)$, where N represents the total number of time discretization steps.[9] This quadratic scaling profile is not limited to solving the deterministic Riccati ODE; when equivalent non-Markovian Volterra integral discretization schemes are applied to Monte Carlo simulation pathways, simulating a single price path containing M time steps similarly demands $\mathcal{O}(M^2)$ floating-point operations.[16]

For high-accuracy pricing algorithms, particularly those seeking to accurately capture the steep volatility skew at very short maturities, the grid density N must frequently be set to significantly large values. Implementations often necessitate up to $N = 20,000$ discrete time steps to ensure minimum stability thresholds are met under extreme parameter regimes.[6] Under a quadratic scaling law, an operation that requires merely milliseconds at low resolutions rapidly devolves into a multi-second computational bottleneck. This exponential decay in performance completely precludes the deployment of the basic fractional Adams method for real-world, real-time model calibration engines, where optimization routines mandate the iterative pricing of thousands of options within strict latency limits.[15]

Numerical Degradation and Stability Limitations

While the fractional Adams method often serves as the foundational benchmark numerical solver in academic literature, it is empirically recognized as notoriously fragile in operational settings. The stability, precision, and convergence behavior of the methodology are exquisitely sensitive to the specified roughness parameter H and the length of the time horizon under consideration. Theoretical convergence proofs for the fractional Adams scheme establish that the maximum absolute error bounded over a local interval scales proportionally to $\mathcal{O}(\Delta^{2-\alpha})$ for values of $x > \epsilon$, and degrades to $\mathcal{O}(\Delta)$ in the immediate vicinity of zero.[9]

However, when this generic fractional solver is specifically applied to the rough Heston characteristic function, the method suffers from profound, non-linear stability decay.[8] For very small Hurst indices approaching zero ($H \to 0$), the singularity inherent in the kernel $(t-s)^{H-1/2}$ becomes severely pronounced. Extensive numerical experiments indicate that while the real component of the fractional Riccati solution generated by the Adams method performs moderately well, the imaginary component degrades rapidly.[6] This specific degradation initiates erratic numerical oscillations, which translate directly into highly erroneous implied volatility surfaces and incorrect option prices unless the errors are artificially suppressed by utilizing extremely dense, computationally prohibitive temporal grids.[6]

Furthermore, the fractional Adams method demonstrates marked instability as the maturity horizon extends. The accumulation of convolution errors over long timeframes forces quantitative researchers to seek alternative stabilization techniques. These workarounds often involve the deployment of multipoint Padé approximations or large-time asymptotic expansions, attempting to manually paste disparate numerical solutions together to cover the required maturity spectrum.[9] Ultimately, owing to the compounded vulnerabilities of $\mathcal{O}(N^2)$ computational overhead, extreme instability in hyper-rough regimes, and poor long-maturity convergence, numerical algorithms predicated purely on the fractional Adams method are broadly deemed inadequate for industrial-grade financial engineering and high-frequency risk management.[13]

Hybrid Simulation Schemes and Fractional Power Series

Recognizing the insurmountable computational and stability barriers imposed by pure fractional convolution schemes, researchers within the quantitative finance community pioneered the development of Hybrid numerical schemes. Although diverse in their specific implementations depending on whether they are applied to ODE solving or Monte Carlo simulation, these methods generally share a unifying mathematical philosophy: they deliberately decouple the numerical treatment of the fractional singularity located near the origin from the smoother, historical integration path, applying distinctly optimized numerical techniques to each respective segment.[18]

Hybrid Schemes for the Fractional Riccati Equation

When tasked with solving the fractional Riccati ODE to derive the characteristic function for Fourier inversion, the standard Adams method effectively forces the numerical solver to repetitively grind through the singularity at zero, instigating the aforementioned instability. An advanced, highly performant alternative is the deployment of a hybrid numerical algorithm that heavily utilizes fractional power series representations. Around the origin $x = 0$, analytical techniques permit the exact representation of the solution to the fractional Riccati equation as a convergent fractional power series.[13]

Because the strict mathematical domain of convergence for this exact power series representation is finite, the Hybrid scheme evaluates the power series analytically only within its rigorously defined stable radius. Once the boundary of convergence is approached, the algorithm seamlessly transitions to an advanced extrapolation technique—frequently a multi-step Richardson-Romberg extrapolation machinery or a finely calibrated standard time-discretization scheme—to approximate the solution for the remainder of the temporal domain.[13]

This hybrid architecture circumvents the massive convolution bottleneck in the highly singular region entirely. The resulting Hybrid scheme is documented to be remarkably fast, exceptionally stable across a wider parameter space, and largely outperforms the fractional Adams method in empirical benchmarks. For short-dated maturities, the Hybrid scheme establishes an evaluation profile where the computational time required to obtain the solution remains virtually constant, independent of the maturity scale, representing a monumental leap in Fourier pricing efficiency.[13]

Hybrid Schemes for Volterra Integration (Bennedsen-Lunde-Pakkanen)

In parallel to Fourier pricing techniques, direct Monte Carlo simulation of non-Markovian rough volatility paths (e.g., simulating fractional Brownian motion directly or evaluating broader Brownian semistationary processes) has benefited tremendously from hybrid methodologies. The Hybrid scheme introduced by Bennedsen, Lunde, and Pakkanen serves as the foundational architectural blueprint for modern fractional Monte Carlo engines.[18]

In this approach, the continuous stochastic Volterra integral is surgically partitioned into two distinct computational components. In the immediate vicinity of the singularity (when the historical integration variable s is exceptionally close to the current forward time t), the fractional kernel is mathematically approximated by a pure, tractably integrable power function. Across this singular boundary, the stochastic integral is evaluated utilizing either exact distributional properties or finely tuned, high-order approximations. Conversely, for the extended historical path where the time differential $(t - s)$ is large and the kernel is demonstrably smooth, the algorithm defaults to a highly optimized standard Riemann-sum discretization protocol.[9]

Rigorous theoretical analysis of this specific Hybrid scheme reveals a massive improvement in weak error convergence rates compared to naive Volterra Euler discretizations. For stochastic integrals driven directly by a fractional Brownian motion defined by a given Hurst index $H \in (0, 1/2)$, applying an exact left-point discretization yields a weak error rate theoretically bounded at the order of $\mathcal{O}(H + 1/2)$. However, deploying a properly weighted Bennedsen-Lunde-Pakkanen Hybrid scheme radically accelerates this convergence profile, achieving an order of $\mathcal{O}((3H + 1/2) \wedge 1)$.[23] This superior, accelerated convergence profile allows practitioners to drastically reduce the required number of discrete time steps while maintaining targeted accuracy bounds, thereby partially mitigating the heavy computational load intrinsic to non-Markovian simulation.

The Hybrid Quadratic Exponential (HQE) Scheme

Building directly upon the theoretical foundations of exact simulation techniques developed for affine processes, the Hybrid Quadratic Exponential (HQE) scheme was engineered specifically to facilitate the robust Monte Carlo simulation of rough affine forward variance models, inherently encompassing the rough Heston model specification.[25] The HQE framework essentially represents a mathematically rigorous, rough adaptation of the highly celebrated Andersen Quadratic Exponential (QE) scheme, which remains the industry standard for simulating the classical Heston model.

Algorithmic Adaptation and Boundary Handling

The standard Euler-Maruyama discretization of square-root processes often encounters catastrophic failures when the variance process closely approaches or breaches the zero boundary, necessitating ad-hoc truncation or reflection protocols that introduce severe biases. Andersen's original QE scheme resolved this by matching moments and mapping the variance updates to either a non-central chi-square distribution or an exponential distribution, depending on the proximity to zero. The HQE scheme mathematically adapts this bifurcated mapping logic to the rough environment.

Specifically, the HQE scheme merges a complex quadratic-exponential approximation protocol for updating the localized state variable with a concurrent Riemann-sum approximation designed to process the historical fractional Volterra integral.[4] Through this synthesis, the HQE scheme elegantly handles the zero-variance boundary condition while respecting the fractional memory of the system. Detailed numerical experiments confirm that the HQE scheme exhibits a highly stable weak rate of convergence of exactly 1 with respect to the total number of time discretization steps, significantly outperforming unweighted fractional Euler schemes in terms of pricing accuracy.[16]

The Persistence of Quadratic Complexity

However, despite its excellent weak convergence properties, its robust handling of extreme boundary conditions at zero variance, and its capacity to mitigate discretization biases, the pure HQE scheme remains structurally and fundamentally bound by the non-Markovian nature of the underlying rough Heston stochastic process. The scheme does not eliminate the historical memory requirement. Because the numerical engine must comprehensively recompute the entirety of the historical integrals—specifically the evaluation of $\int_0^t K(t-s)\nu\sqrt{V_s}dW_s$ —at every single forward time step, its theoretical asymptotic computational cost remains quadratic.[16]

Consequently, the operational complexity scales rigidly as $\mathcal{O}(M^2)$ per generated sample path, where M denotes the number of time steps. Thus, while the HQE scheme is definitively numerically superior, far more stable, and less biased than naive Euler integration routines, executing pure HQE without integrating advanced Markovian lifting enhancements still forces the computational infrastructure to endure the exceptionally heavy operational burden of full historical discrete convolutions.[16]

Markovian Approximations and the Lifted Heston Architecture

The most definitive structural breakthrough in resolving the crippling computational complexities native to the rough Heston model is the introduction and rigorous application of

Markovian approximations, most prominently formalized and popularized as the "Lifted Heston" model framework.[1] This specific mathematical methodology does not merely attempt to accelerate the integration of the fractional kernel; rather, it completely addresses and eliminates the root cause of the $\mathcal{O}(N^2)$ complexity overhead: the infinite, long-memory nature of the fractional Volterra kernel itself.

Sum-of-Exponentials Kernel Approximation and Bernstein's Theorem

The defining mathematical feature of the fractional kernel utilized within the rough Heston model—specifically formalized as $K(t) = (t - s)^{H-1/2}/\Gamma(H + 1/2)$ —is that it qualifies as a completely monotone function over the interval $(0, \infty)$.[27] According to Bernstein's fundamental theorem on monotone functions, any such completely monotone function can be exactly represented as a Laplace transform of a non-negative Borel measure. Translated into practical numerical applications, this mathematical property dictates that the singular fractional kernel can be approximated with an arbitrarily high degree of precision by deploying a finite, weighted sum of exponential functions:

$$K_N(t) \approx \sum_{i=1}^{N} c_i e^{-x_i t}$$

In this approximation, the specific expansion weights c_i and the mean-reversion speed parameters x_i are carefully determined. Practitioners typically calculate these coefficients via sophisticated Gaussian quadrature rules or through direct numerical optimization techniques strictly targeted at minimizing the L^1 norm error between the true fractional kernel and the exponential sum over the specific time interval of interest.[1]

The N-Dimensional Lifted State Space

By systematically substituting the singular, non-Markovian fractional kernel with this finite sum of smooth exponentials, the underlying stochastic Volterra equation is mathematically "lifted" from its infinite-dimensional, path-dependent state into a finite, high-dimensional, purely Markovian dynamical system. The singular historical variance process V_t is entirely replaced by an aggregated sum of N auxiliary state processes (frequently referred to in the literature as specific model factors). Each individual auxiliary factor evolves independently as a standard mean-reverting stochastic differential equation, with all factors continuously sharing a single, common Brownian diffusion driver.[27] The resultant system is structured as follows:

$$dU_t^{(i)} = (-x_i U_t^{(i)} - \lambda V_t)dt + \nu\sqrt{V_t}dW_t$$

$$V_t = g_0(t) + \sum_{i=1}^{N} c_i U_t^{(i)}$$

This specific factor parametrization rigorously guarantees that as the total number of auxiliary factors approaches infinity ($N \to \infty$), the Lifted Heston multidimensional model converges weakly to the exact, mathematically pure rough Heston model characterized by the target Hurst index H.[29] Remarkably, extensive empirical testing reveals that exceptional numerical stability and high calibration accuracy can be achieved with a surprisingly small number of discrete factors. Depending on the required precision tolerances and the length of the pricing maturity, utilizing $N \in \{2, 10, 20\}$ factors is generally sufficient for industrial-grade applications.[16]

Structural Memory Dependency: Volterra vs. Markovian Lifts

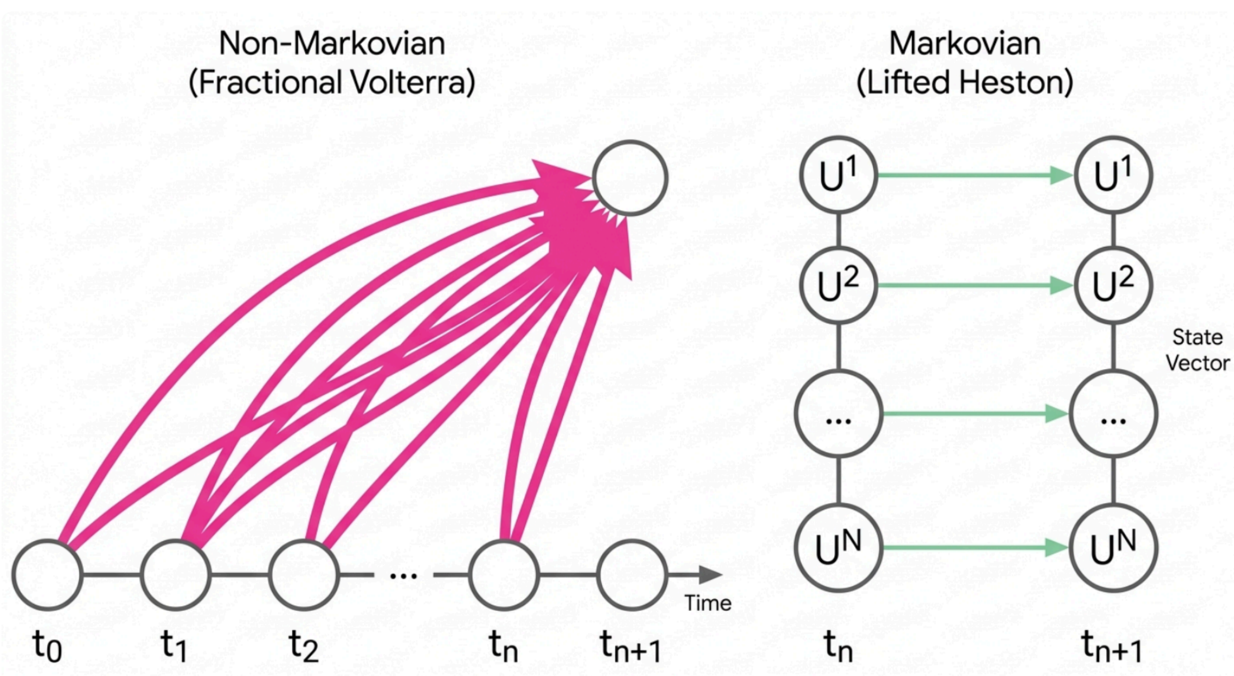

The non-Markovian fractional models require scanning the entire historical path (t_0 to t_n) to compute the next step, causing a memory bottleneck. The Markovian Lifted Heston approximates this history into an N-dimensional state vector that propagates strictly step-by-step, unlocking linear complexity.

Linearization of Complexity and Advanced Weak Schemes

The operational consequence of executing the Lifted Heston Markovian approximation is profoundly transformative for computational finance workflows. Because each of the N auxiliary processes $U_t^{(i)}$ is mathematically defined as a purely Markovian diffusion, its physical state at forward time t_{k+1} is completely determined by its immediate state at time t_k. The model comprehensively sheds the entirety of its historical path dependency.
Consequently, the computational processing cost required to simulate a single Monte Carlo path spanning M time steps precipitously drops from the restrictive $\mathcal{O}(M^2)$ bound down to a highly tractable $\mathcal{O}(N \times M)$ operation.[1] Given that N acts as a fixed, small constant across the simulation, the overall algorithmic complexity scales linearly with the number of time steps. In demanding practical applications, the deployment of the Lifted Heston Markovian approximation truncates the computation time required to generate a single implied volatility surface by massive factors when directly compared to pure rough Heston Volterra integration schemes.[31]
Furthermore, because the resulting multidimensional architecture is strictly Markovian, it natively supports highly advanced, fast simulation techniques historically reserved for standard diffusions. For instance, clipped Euler-Maruyama schemes can be deployed to strictly ensure the non-negativity of the variance process while maintaining linear time complexity.[32]
Additionally, newer numerical innovations, such as the Integrated Volterra Implicit (iVi) scheme based on Inverse Gaussian distributions, have been explicitly developed for these Markovian structures. The iVi scheme elegantly handles singularities by relying solely on integrated kernel quantities, preserving the non-decreasing property of integrated processes, and establishing robust weak convergence with exceptionally few discretization steps.[32]

Comparative Profiling of Algorithmic Trade-offs

Evaluating these numerical schemes strictly through the isolated lens of asymptotic time complexity provides an incomplete assessment; the comparative numerical stability, the established rates of weak convergence, and the scheme's behavioral properties at extreme boundaries (especially near zero variance and in hyper-rough $H \ll 1/2$ regimes) are equally critical criteria for industrial deployment.

Performance Metric	**Fractional Adams (Pure)**	**Hybrid Schemes (HQE / Power Series)**	**Markovian Approx. (Lifted Heston)**

Asymptotic Complexity	$\mathcal{O}(M^2)$	$\mathcal{O}(M^2)$ (MC) / Constant (Riccati)	$\mathcal{O}(N \times M)$ (Strictly Linear)
Theoretical Weak Convergence Rate	$\mathcal{O}(\Delta)$ (severely deteriorates near singularity)	$\mathcal{O}(1)$ (Specifically for HQE framework)	Ranging from $\mathcal{O}(\Delta)$ up to $\mathcal{O}(\Delta^2)$
Stability Under $H \to 0$ Regime	Highly unstable, generates severe oscillatory pricing errors	Very stable due to explicit mathematical singularity handling	Exceptionally stable, scaling robustly with increased N factors
Variance Positivity Guarantee	None natively embedded in the mathematical structure	High (via Quadratic Exponential distributional mapping)	High (via Clipped Euler techniques or Inverse Gaussian weak schemes)
Suitability for Path-Dependent Pricing	Prohibitive due to insurmountable $\mathcal{O}(M^2)$ cost	Feasible but remains operationally slow	Highly efficient, enabling standard Longstaff-Schwartz regressions

Stability Under Extreme Roughness Dynamics

The practical utility of the rough Heston model is uniquely maximized under the exact market conditions where it is most mathematically difficult to simulate: at extraordinarily low Hurst indices (specifically $H \in [0.05, 0.15]$). These highly singular parameters are absolutely necessary to successfully calibrate the model to the notoriously steep, short-term implied volatility skews chronically exhibited by the SPX and VIX options markets.[2]

Under these extreme, hyper-rough parameters, the standard Fractional Adams method exhibits severe and rapid numerical decay. Because the integration kernel $(t-s)^{H-1/2}$ violently approaches a non-integrable singularity as $H \to 0$, the standard polynomial interpolation mathematics underlying the Adams corrector algorithm completely fails to capture the aggressive local dynamics, resulting in pricing breakdown.[6]

The Hybrid HQE scheme bypasses this specific failure mode by analytically pre-computing the exact integral over the local singular interval and applying the robust Quadratic Exponential mapping to update the state variable smoothly across the boundary.[12] However, perhaps most remarkably, empirical back-tests conducted on advanced Markovian approximations (specifically referencing the iVi scheme for Volterra Heston structures) demonstrate that the mathematical stability and convergence rates of Lifted models actually *improve* as the assigned Hurst index structurally decreases and aggressively approaches the theoretical limit of $-1/2$.[27] This counter-intuitive resilience under extreme duress makes Markovian lifting the definitively superior mathematical framework for navigating deep-rough, hyper-singular volatility regimes.

Hardware-Level Execution: GPU Parallelization Dynamics

In modern quantitative finance and industrial risk management, the theoretical asymptotic complexity of an algorithmic model is intrinsically and permanently tied to its hardware execution profile. Massive computational tasks, such as real-time option pricing matrices and global volatility surface calibration via extensive Monte Carlo simulation or high-dimensional PDE solving, represent highly parallel workloads. Consequently, they are prime candidates for execution on Graphics Processing Units (GPUs) or specialized Tensor Processing Units (TPUs).[33] GPUs operate fundamentally on a Single Instruction, Multiple Threads (SIMT) architecture. To achieve maximum computational throughput and justify the hardware investment, an algorithm must relentlessly prioritize high compute-to-memory-access ratios. Furthermore, it must strictly minimize warp divergence (a hardware condition where adjacent threads are forced to follow different logical branching paths) and ensure perfectly coalesced memory access (where adjacent computational threads concurrently read from adjacent physical memory addresses).[35] When subjected to these unforgiving hardware constraints, the three primary numerical schemes exhibit vastly different compatibility profiles.

The Fractional Adams Bottleneck: Memory Bandwidth Bound

The Fractional Adams method, and by theoretical extension any pure non-Markovian Volterra integration scheme, is catastrophically unsuited for large-scale GPU acceleration architectures. The fundamental $\mathcal{O}(M^2)$ mathematical complexity translates directly into a paralyzing memory bandwidth bottleneck at the silicon level. To compute the forward step t_{k+1}, a designated CUDA thread must sequentially read the entire continuous historical array of stored states from t_0 through t_k.

As the simulation time index k grows throughout the path, the historical data array rapidly exceeds the limited capacity of the ultra-fast L1/Shared memory caches (which typically range from a mere 64KB to 128KB per streaming multiprocessor on advanced enterprise architectures like the NVIDIA Tesla V100).[33] Once the historical data is inevitably forced out into the much slower, high-latency Global Memory banks, the algorithm becomes totally memory-bound. Even if the historical read accesses are perfectly coalesced across the thousands of executing Monte Carlo paths, the sheer volumetric requirement of data transfer thoroughly saturates the GPU's memory bus. This saturation leaves the thousands of available Arithmetic Logic Units (ALUs) idling, consuming power while waiting for data retrieval. Although highly complex hybrid implementations utilizing asynchronous CPU-GPU coordination can attempt to intelligently partition the workload, the fundamental mathematical heredity (infinite memory) of the fractional differentiation operator consistently chokes parallel processing efficiency.[35]

The HQE Scheme: Divergence and Latency Truncation

The Hybrid Quadratic Exponential (HQE) scheme performs significantly better on SIMT GPU architectures than the standard Fractional Adams protocol. This improvement primarily arises because the integrated QE component provides exact, large-step distributional transitions that structurally limit the total number of required time steps M needed to reach maturity.[25] However, when evaluating the historical integral component of the path, the HQE scheme still mandates the scanning of past states, thereby inheriting a diluted form of the memory bandwidth bottleneck.

Furthermore, the HQE algorithm relies heavily on internal conditional branch operations. The code must dynamically switch between executing the quadratic formula and the exponential formula depending entirely on whether the instantaneous variance level crosses a specific threshold. On a GPU, dynamic conditional branching within a warp induces warp divergence. When divergence occurs, threads within the same 32-thread execution warp evaluate different conditions and must execute their respective instructional branches sequentially rather than in parallel, effectively halving the hardware efficiency. While exceptionally careful, low-level CUDA coding can attempt to mitigate these divergence impacts, the residual historical memory footprint ultimately remains a firm limiting factor on total throughput.[12]

The Lifted Heston Triumph: Compute Bound and Coalesced Access

The structural architecture of the Markovian approximation fundamentally, and almost perfectly, aligns with the operational requirements of CUDA hardware.[34] By mathematically transforming the singular infinite-memory system into exactly N parallel, memoryless state variables, the Lifted Heston model guarantees that to step forward in time, an executing thread only ever needs to access the N variables from the immediately preceding temporal step. For a standard configuration of $N = 20$ factors, the entire active state vector fits exceptionally comfortably within the fastest available tier of GPU registers or L1 cache.[28] Under this architecture, there is zero requirement to query the slow global memory banks for historical path data. The algorithmic execution completely flips from being memory-bound to

being purely compute-bound, allowing the GPU to fully unleash its thousands of ALUs at maximum clock rates.
Furthermore, because all executing Monte Carlo paths apply the exact same clipped Euler or weak numerical updates continuously without triggering complex dynamic branching thresholds, warp divergence is virtually eliminated across the multiprocessor.[32] Memory write accesses for updating the state vectors across hundreds of thousands of concurrent paths can be perfectly coalesced into highly efficient, contiguous 128-byte transactions.[38]
This profound architectural harmony directly enables quantitative researchers and risk managers to leverage high-performance computing clusters (such as arrays of NVIDIA Tesla V100S GPUs on supercomputers like DelftBlue) to simulate millions of complex stochastic paths in small fractions of a second. This unprecedented speed not only enables the real-time calibration of the rough Heston model against live market feeds but also unlocks the application of massive machine learning tools. Advanced methodologies like Time-Stepping Deep Gradient Flow networks or complex Neural PDE Solvers intrinsically mandate the existence of Markovian state spaces to function efficiently, further solidifying the necessity of the lifting approach.[33]

Conclusions and Forward Outlook

The relentless pursuit of an optimal, industrially scalable numerical scheme for the rough Heston model necessitates navigating an extraordinarily complex tradeoff landscape between theoretical mathematical exactness, numerical stability boundaries, and raw hardware operational feasibility. The characteristic function's immutable dependence on the non-linear fractional Riccati equation, coupled inextricably with the non-Markovian nature of the fractional Volterra simulated paths, originally established a massive, seemingly intractable computational hurdle for early volatility practitioners.
The foundational fractional Adams-Bashforth-Moulton method, while conceptually straightforward and theoretically robust under standard conditions, is fundamentally and fatally hindered by its $\mathcal{O}(M^2)$ asymptotic complexity. This scaling failure is compounded by severe mathematical stability degradation precisely in the deep-rough, singular regimes ($H \to 0$) demanded by empirical market volatility data. Its uncompromising reliance on total path history results in rigid algorithms that are perpetually memory-bandwidth bound and inherently hostile to modern GPU SIMT parallel architectures.
Hybrid schemes, such as the highly refined HQE scheme for Monte Carlo simulation and fractional power-series methods for ODE Riccati solving, represent a highly robust, mathematically elegant middle ground. By explicitly treating the integration singularity near zero through exact methods and leveraging stable numerical transitions for smoother historical regions, these Hybrid schemes drastically improve weak convergence rates and extend numerical stability. However, without structural, dimensional modifications, Volterra-based hybrid schemes still stubbornly incur the quadratic scaling costs of historical lookup operations.
The most profound and practically impactful advancement in this highly specialized domain remains the Markovian approximation methodology, perfectly epitomized by the Lifted Heston

model architecture. By rigorously approximating the completely monotone fractional kernel with an optimized sum of exponentials, the methodology elegantly, and legally within the bounds of weak convergence theorems, collapses the historical path dependency into a finite, highly manageable state space. This mathematical breakthrough single-handedly shifts the algorithmic complexity from an unworkable quadratic boundary down to a strictly linear ($\mathcal{O}(N \times M)$) scaling profile. It completely transforms the simulation engine into a compute-bound operation perfectly suited for massive GPU parallelization, effortlessly boasting coalesced memory access architectures and minimizing detrimental warp divergence.

As the quantitative finance industry continues its inevitable trajectory toward demanding higher-frequency model calibration cycles and the pricing of increasingly complex, deeply path-dependent exotic structures, the absolute necessity for high-throughput, GPU-accelerated algorithms renders the Markovian approximation not merely the most efficient operational choice, but the mathematically indispensable standard for deploying the rough Heston model. Looking forward, emerging quantitative innovations—ranging from signature-based pricing frameworks that map volatility as linear functionals, to advanced deep neural network PDE solvers utilizing reservoir computing—will undoubtedly continue to build heavily upon this stable Markovian foundation. The progression of these technologies will likely result in the complete obsolescence of direct historical convolution methods within elite financial production environments, establishing lifted Markovian dynamics as the permanent architecture for rough volatility modeling.

Works cited

1. efficient option pricing in the rough heston model using weak ..., accessed May 27, 2026, https://www.wias-berlin.de/people/bayerc/files/bayer_breneis23b_arxiv.pdf
2. Fast reliable pricing and calibration of the rough Heston model - arXiv, accessed May 27, 2026, https://arxiv.org/html/2508.15080
3. Perfect hedging in rough Heston models - Project Euclid, accessed May 27, 2026, https://projecteuclid.org/journals/annals-of-applied-probability/volume-28/issue-6/Perfect-hedging-in-rough-Heston-models/10.1214/18-AAP1408.pdf
4. Rough Heston model as the scaling limit of bivariate cumulative heavy-tailed INAR processes: Weak-error bounds and option pricing - arXiv, accessed May 27, 2026, https://arxiv.org/html/2503.18259v6
5. Multifactor Approximation of Rough Volatility Models - SIAM Publications Library, accessed May 27, 2026, https://epubs.siam.org/doi/pdf/10.1137/18M1170236?download=true
6. Fractional Riccati equation and its applications to rough Heston model - Preprints.org, accessed May 27, 2026, https://www.preprints.org/manuscript/202002.0311/v1/download
7. arXiv:1711.03078v4 [math.PR] 11 Nov 2023, accessed May 27, 2026, https://arxiv.org/pdf/1711.03078

8. Solving Parametric Fractional Differential Equations Arising from the Rough Heston Model Using Quasi-Linearization and Spectral Collocation | SIAM Journal on Financial Mathematics, accessed May 27, 2026, https://epubs.siam.org/doi/10.1137/19M1269324
9. Fractional Riccati Equation and Its Applications to Rough Heston ..., accessed May 27, 2026, https://www.mdpi.com/2073-8994/12/6/959
10. Rough Volatility in Financial Mathematics - SIAM, accessed May 27, 2026, https://www.siam.org/publications/siam-news/articles/rough-volatility-in-financial-mathematics/
11. Pricing in the Heston Model and Its Rough Variation - DiVA portal, accessed May 27, 2026, https://www.diva-portal.org/smash/get/diva2:1800508/FULLTEXT01.pdf
12. Pricing in affine forward variance models - Baruch MFE Program, accessed May 27, 2026, https://mfe.baruch.cuny.edu/wp-content/uploads/2024/03/AFVPricingDevelopmentsQuantminds2022.pdf
13. Fast Hybrid Schemes for Fractional Riccati Equations (Rough is not ..., accessed May 27, 2026, https://arxiv.org/abs/1805.12587
14. Asymptotics of the Rough Heston Model - University of Cape Town, accessed May 27, 2026, https://open.uct.ac.za/bitstreams/f6191e57-6d84-44f6-aa55-0c2803d04253/download
15. Option pricing for rough Heston model using numerical methods, accessed May 27, 2026, http://psasir.upm.edu.my/id/eprint/104713/
16. Efficient option pricing in the rough Heston model using weak simulation schemes - Taylor & Francis, accessed May 27, 2026, https://www.tandfonline.com/doi/full/10.1080/14697688.2024.2391523
17. Deep calibration with random grids - arXiv, accessed May 27, 2026, https://arxiv.org/html/2306.11061v2
18. [1507.03004] Hybrid scheme for Brownian semistationary processes - arXiv, accessed May 27, 2026, https://arxiv.org/abs/1507.03004
19. decoupling the short- and long-term behavior of stochastic volatility - College of Business, accessed May 27, 2026, https://www.cb.cityu.edu.hk/ef/doc/2016%20Sofie/Papers/153_Bennedsen_Volatility%20modelling%20decoupling%20the%20short-%20and%20long-term%20behavior%20of%20stochastic%20volatility.pdf
20. Fast Hybrid Schemes for Fractional Riccati Equations (Rough is not so Tough) - arXiv, accessed May 27, 2026, https://arxiv.org/pdf/1805.12587
21. Hybrid scheme for Brownian semistationary processes - Aarhus, accessed May 27, 2026, https://pure.au.dk/portal/en/publications/hybrid-scheme-for-brownian-semistationary-processes/
22. Markovian Approximation of the Rough Bergomi Model for Monte Carlo Option Pricing - Monash University, accessed May 27, 2026, https://www.monash.edu/__data/assets/pdf_file/0016/2320414/WP_CQFIS_2020_2.pdf

23. Weak Error Rates of Numerical Schemes for Rough Volatility | SIAM Journal on Financial Mathematics, accessed May 27, 2026, https://epubs.siam.org/doi/10.1137/22M1485760
24. Weak Error Rates of Numerical Schemes for Rough Volatility | Request PDF - ResearchGate, accessed May 27, 2026, https://www.researchgate.net/publication/370891225_Weak_Error_Rates_of_Numerical_Schemes_for_Rough_Volatility
25. Efficient simulation of affine forward variance models - Risk.net, accessed May 27, 2026, https://www.risk.net/media/download/1075381/download
26. Rough volatility moves to exotic frontiers - Risk.net, accessed May 27, 2026, https://www.risk.net/cutting-edge/views/7928716/rough-volatility-moves-to-exotic-frontiers
27. Lifting the Heston model | Request PDF - ResearchGate, accessed May 27, 2026, https://www.researchgate.net/publication/334306304_Lifting_the_Heston_model
28. Fast reliable pricing and calibration of the rough Heston model - arXiv, accessed May 27, 2026, https://arxiv.org/html/2508.15080v1
29. Lifted Heston Model: Efficient Monte Carlo Simulation with Large Time Steps - arXiv, accessed May 27, 2026, https://arxiv.org/html/2510.08805v1
30. Rough Hawkes Heston Model for Volatility - Scribd, accessed May 27, 2026, https://www.scribd.com/document/885064796/simone-scotti-the-rough-hawkes-heston-stochastic-volatility-model-1
31. Comparing Rough Volatility Models - TU Delft Repository, accessed May 27, 2026, https://repository.tudelft.nl/file/File_cd0ef62c-e82f-466a-bf60-91445385b07a
32. Simulating integrated Volterra square-root processes and Volterra Heston models via Inverse Gaussian - arXiv, accessed May 27, 2026, https://arxiv.org/html/2504.19885v1
33. A time-stepping deep gradient flow method for option pricingΣin (rough) diffusion models, accessed May 27, 2026, https://arxiv.org/html/2403.00746v1
34. A time-stepping deep gradient flow method for option pricing in rough diffusion models - TU Delft Research Portal, accessed May 27, 2026, https://research.tudelft.nl/files/262757347/A_time-stepping_deep_gradient_flow_method_for_option_pricing_in_rough_diffusion_models.pdf
35. Fractional Riccati Equation and Its Applications to Rough Heston Model Using Numerical Methods - ResearchGate, accessed May 27, 2026, https://www.researchgate.net/publication/341944304_Fractional_Riccati_Equation_and_Its_Applications_to_Rough_Heston_Model_Using_Numerical_Methods
36. Unsupervised Learning-based Calibration Scheme for Rough Volatility Models - arXiv, accessed May 27, 2026, https://arxiv.org/pdf/2412.02135
37. [PDF] Hybrid GPU–CPU Efficient Implementation of a Parallel, accessed May 27, 2026, https://www.semanticscholar.org/paper/Hybrid-GPU%E2%80%93CPU-Efficient-Implementation-of-a-for-a-Tverdyi-Parovik/60f23113822d0859eaf00988193f145729034477
38. (PDF) On the Discrete-Time Simulation of the Rough Heston Model - ResearchGate, accessed May 27, 2026,

https://www.researchgate.net/publication/368355471_On_the_Discrete-Time_Simulation_of_the_Rough_Heston_Model
39. Lech Grzelak Doctor of Philosophy Professor (Associate) at Utrecht University - ResearchGate, accessed May 27, 2026, https://www.researchgate.net/profile/Lech-Grzelak
40. Random Neural Networks for Rough Volatility - PMC - NIH, accessed May 27, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC12967488/